

# The Docker Revolution

by @twombh February 2017

# Vertical versus Horizontal



# How to find a web page

## Vertically

Step 1: Get a very, very, very big computer.

Step 2: Get a very, very, very fast Internet connection.

Step 3: Load every page on the Internet.



## Horizontally

Step 1: Get lots of computers to scrape the Internet. Yesterday.

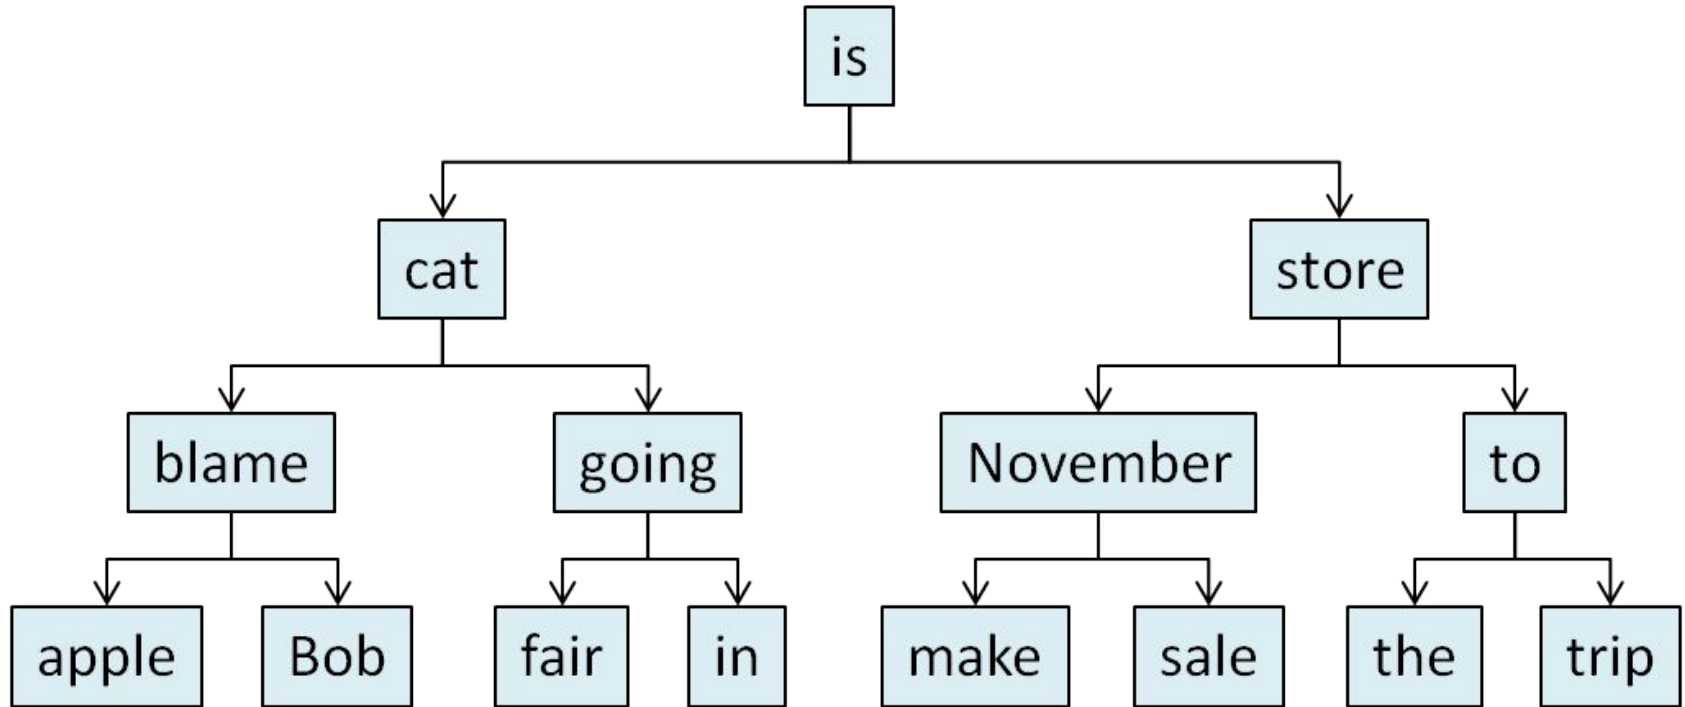
Step 2: Store just the text using something like Balanced Binary Trees.

Step 3: Search the Internet in RAM.

# The limits of physics



# Balanced Binary Search Tree



# Concurrency versus Parallelism





# Interruptible versus Independent




Walking on the street is **interruptible**.  
But answering your phone is **not independent**.



Putting out a fire is **independent**.  
But **not interruptible**.

# Race Conditions





# Handling lots of web requests



- \* Independent, but not interruptible.  
Eg; failed email delivery messages
- \* Racks with 256GB RAM and 16 cores :)
- \* VMs are programmable, which saves labour and money.
- \* **Arbitrary scaling** and parallelism mean we can forget about physics.
- \* ‘Elastic computing’ saves money.


# Independence As A Service



- \* Heroku was a natural evolution of shared hosting. But not using VMs.
- \* Less access to servers (no FTP, no permanent storage, etc) and strict requirements (git, all external services must be run separately).
- \* But it could be easily and arbitrarily scaled.

# Formalisation of Parallelisation

- \* LXC: Linux containers.
- \* Google and other large compute companies.
- \* Doesn't virtualise CPU/RAM it restricts access to them.
- \* More like `init` process than a VM.
- \* PHP, Python, Ruby, Java: all run on Linux.



# Solomon Hykes and dotCloud



- \* dotCloud were essentially a Heroku competitor.
- \* Hykes was using LXC.
- \* Docker was borne from the tools Hykes was building at dotCloud.
- \* March 2013.
- \* He was simply open sourcing a way of parallelisation that had been around for a few years.

# The 12 Factor App Manifesto

**I. Codebase** One codebase tracked in revision control, many deploys.

**II. Dependencies** Explicitly declare and isolate dependencies.

**III. Config** Store config in the environment.

**IV. Backing services** Treat backing services as attached resources.

**V. Build, release, run** Strictly separate build and run stages.

**VI. Processes** Execute the app as one or more stateless processes.

**VII. Port binding** Export services via port binding.

**VIII. Concurrency** Scale out via the process model.


**IX. Disposability** Maximize robustness with fast startup and graceful shutdown.

**X. Dev/prod parity** Keep development, staging, and production as similar as possible.

**XI. Logs** Treat logs as event streams.

**XII. Admin processes** Run admin/management tasks as one-off processes.

# Codebase in Revision Control



- \* Backup, querying history, collaborating (like a mutex).
- \* Moving code between live, distributed, concurrent, containers is quicker.
- \* Prevents hotfixing culture.
- \* Better for onboarding new devs.
- \* ``docker pull/push``.

# Explicitly Declare Dependencies

- \* Code does not exist in isolation.  
Depends on DB protocol versions  
etc.
- \* Independence must carefully  
acknowledge dependence.
- \* `docker build`

```
package.json
1 {
2   "name": "npm-check",
3   "version": "0.0.1",
4   "description": "Comparing NPM (dev)Dependencies",
5   "dependencies": {
6     "bootstrap": "^3.3.4",
7     "browserify": "^9.0.8",
8     "chart.js": "^1.0.2",
9     "font-awesome": "^4.3.0",
10    "react": "^0.13.2",
11    "react-bootstrap": "^0.21.0",
12    "react-chartjs": "^0.6.0",
13    "react-dropzone": "^1.0.1",
14    "react-router": "^0.13.3",
15    "react-tools": "^0.13.2",
16    "reactify": "^1.1.0",
17    "reflux": "^0.2.7",
18    "superagent": "^1.2.0",
19    "underscore": "^1.8.3",
20    "watchify": "^3.1.2"
21  },
22  "devDependencies": {
23    "grunt": "^0.4.5",
24    "grunt-contrib-clean": "^0.6.0",
25    "grunt-contrib-copy": "^0.8.0",
26    "grunt-contrib-less": "^1.0.1",
27    "grunt-contrib-watch": "^0.6.1"
```




# Store Config in the Environment



- \* Better for testing. Forces you to decouple code.
- \* Single Responsibility Principle.
- \* The 4 cornerstone environments: developers/testing/workers/production.
- \* Hot-swapping dead databases.
- \* ``docker run --env``

# Testing Gives You Half of Docker

- \* Dependencies, Config, Backing Services, Separated Build/Run, Disposability.
- \* Testing requires a fundamental shift in how you write code.
- \* Similar to Docker, it requires fundamentally changing your workflow and the way you think.



# The Utopian Vision



- \* Docker isn't just a thing to put apps in.
- \* It aims to improve the entire world of creating, testing, collaborating, deploying, updating and scaling apps.

# Democratisation of the Internet

- \* If an app can be completely made independent then there is no need for the company surrounding the app.
- \* Apps can be deployed directly from Github with no devops.
- \* Imagine if all the essential Internet services such as search and social media were funded like Wikipedia.



# Where to Start?

Checkout **Awesome Docker** on Github

<https://github.com/veggemonk/awesome-docker>